Clustering Algorithms with 2018 BRFSS Data

By Heather Knudson

Behavioral Risk Factor Surveillance System

- Yearly telephone surveys in all 50 states, D.C., & 3 territories conducted by the CDC
- More than 400,000 interviews each year
- Questions on health, behavior, demographics

BRFSS™

Data Cleaning

Data Cleaning

1. Removed columns with > 20% of data missing
2. For columns with < 20% missing, imputed code ‘1000 to stand in for missing values
3. For continuous variables, imputed with the mean
4. Removed columns that were copies of other columns, favoring those with the most info
5. Re-coded all columns so the first code was 0, not 1
6. Re-coded columns that were originally in ‘backwards’ order

Total: 437,436 observations & 68 features

Research Question

How can 2018 BRFSS respondents be segmented based on their demographics, behaviors, and health outcomes?

Descriptive Statistics

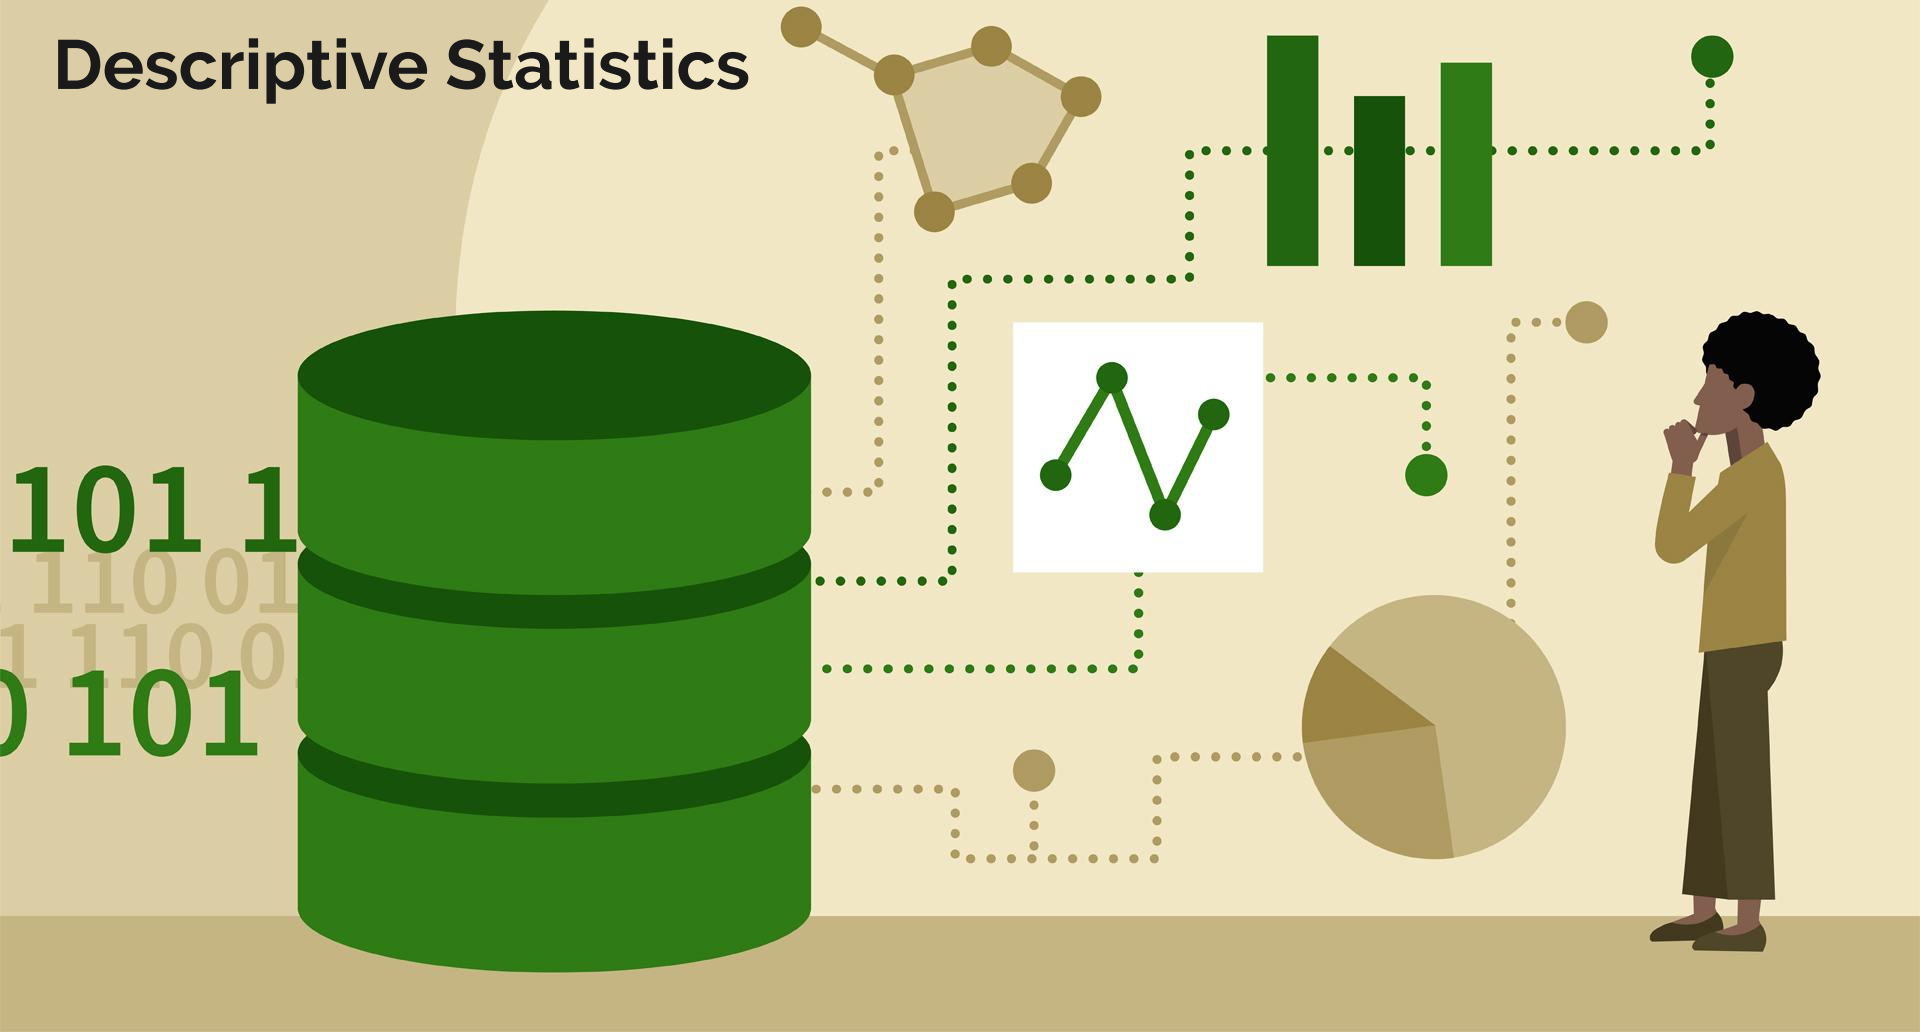

Most respondents lived in NY, followed by MN

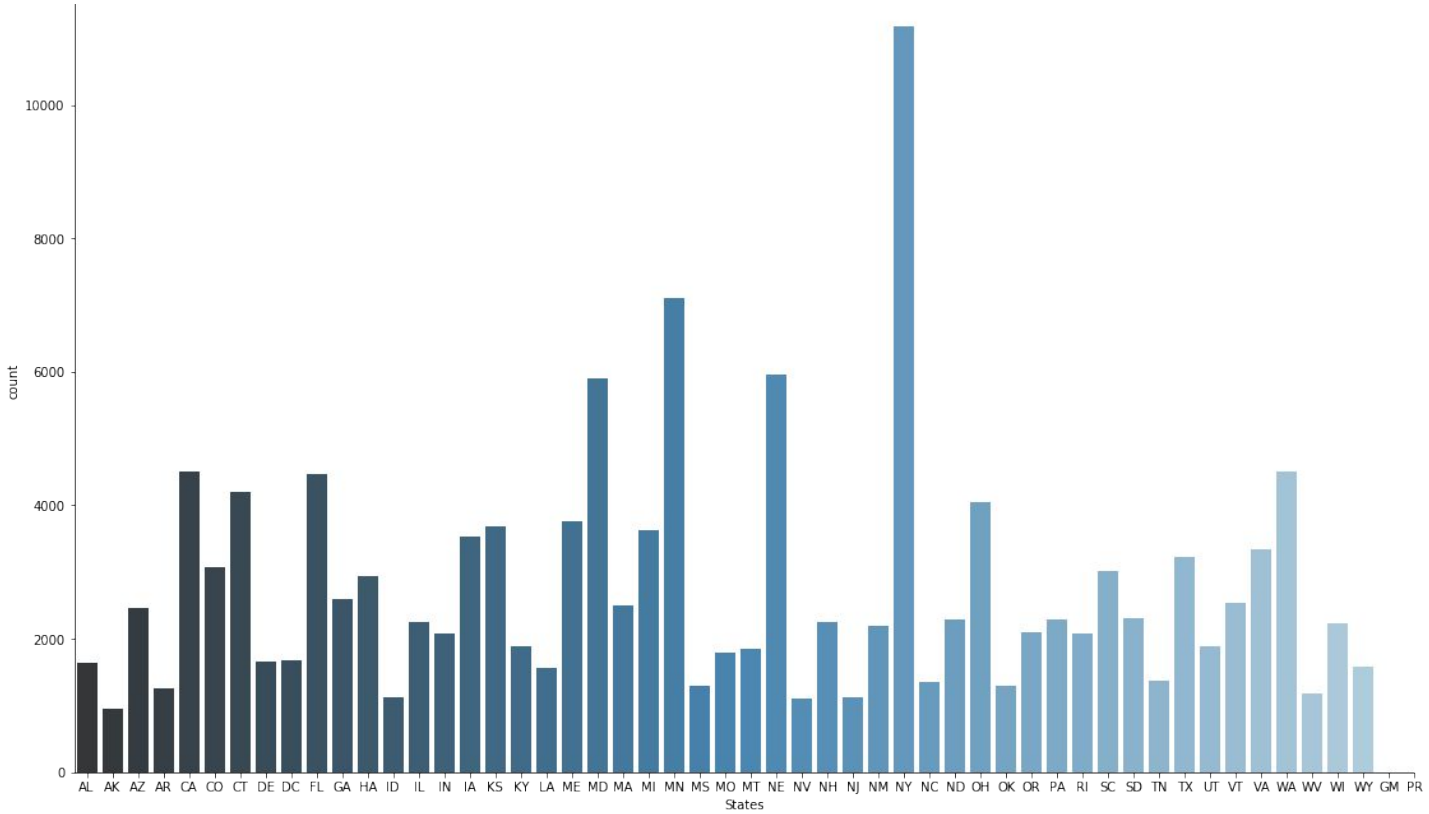

Most respondents were in their 50s & 60s

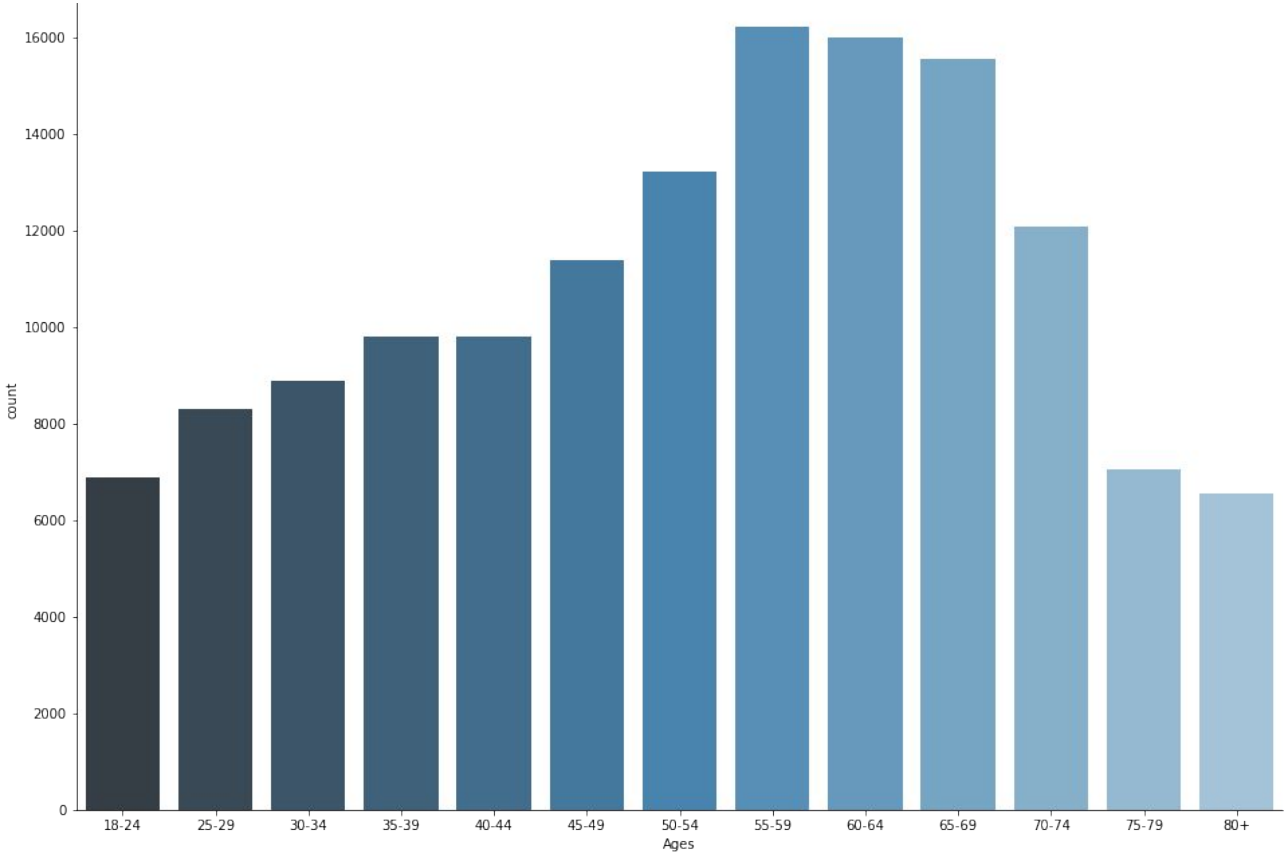

Most respondents made more than $75,000

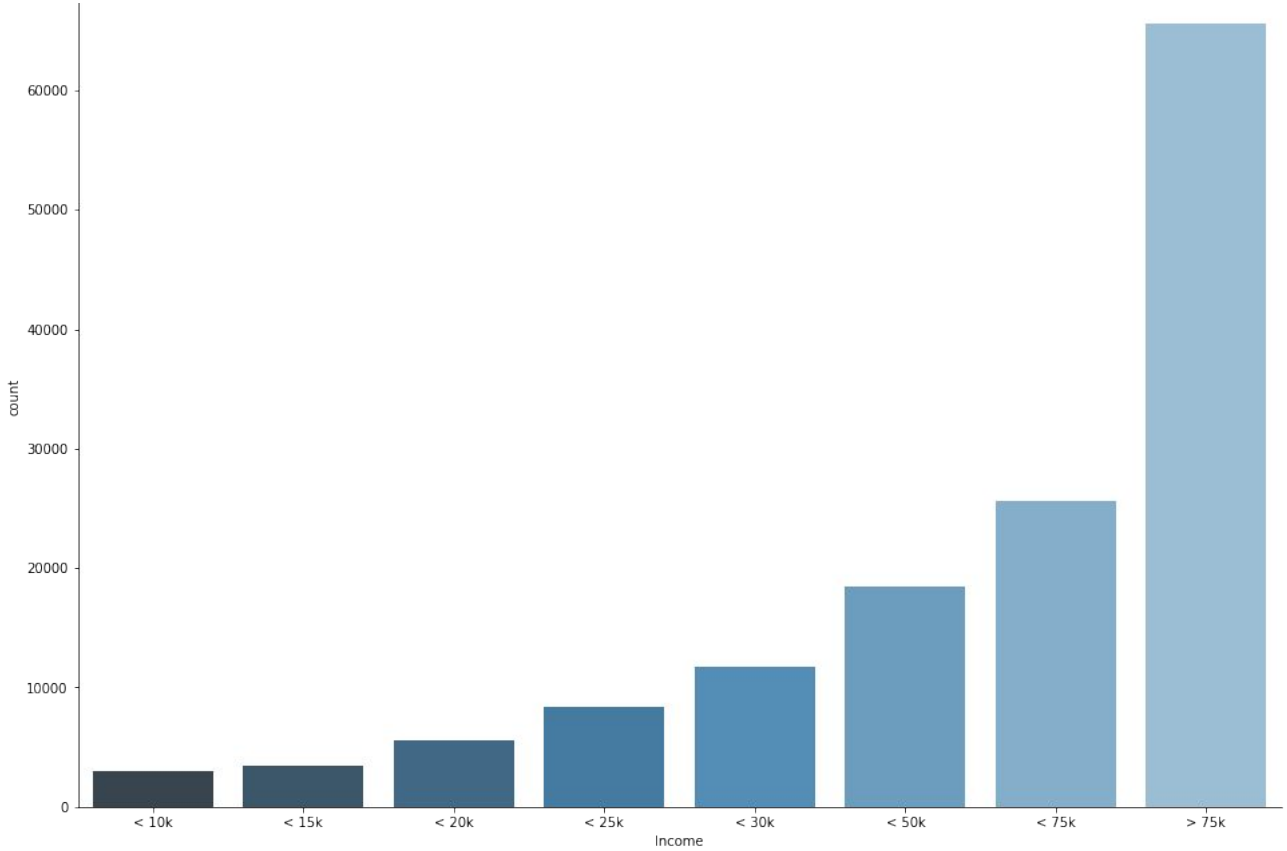

Most respondents had a college degree

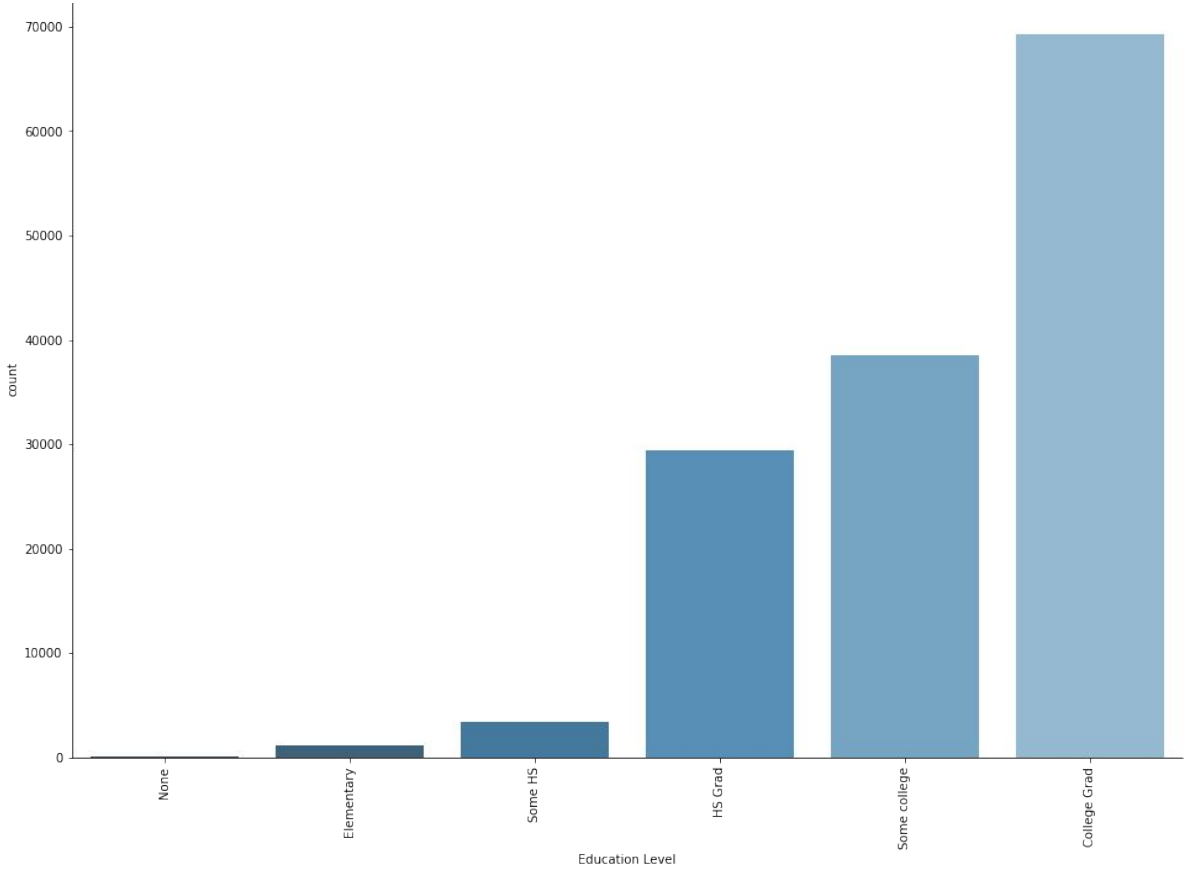

Most respondents were in good physical health

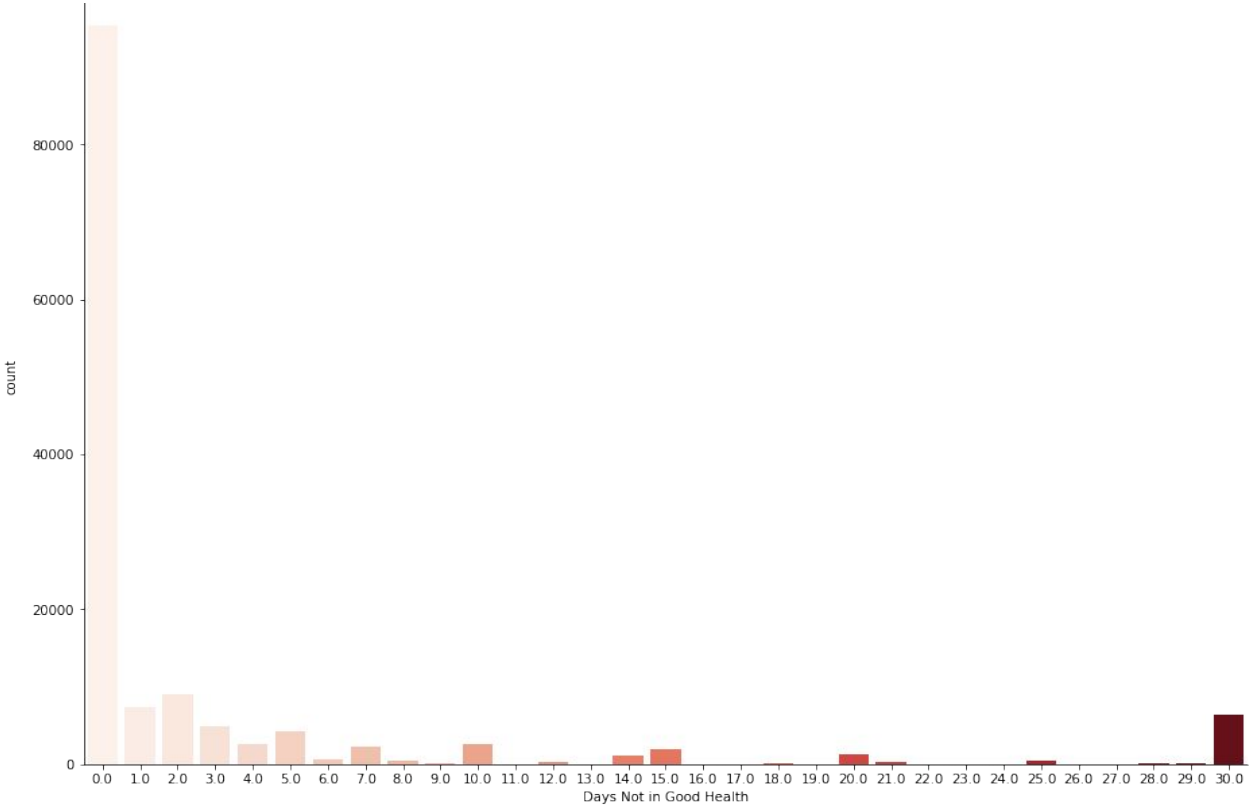

Most respondents were in good mental health

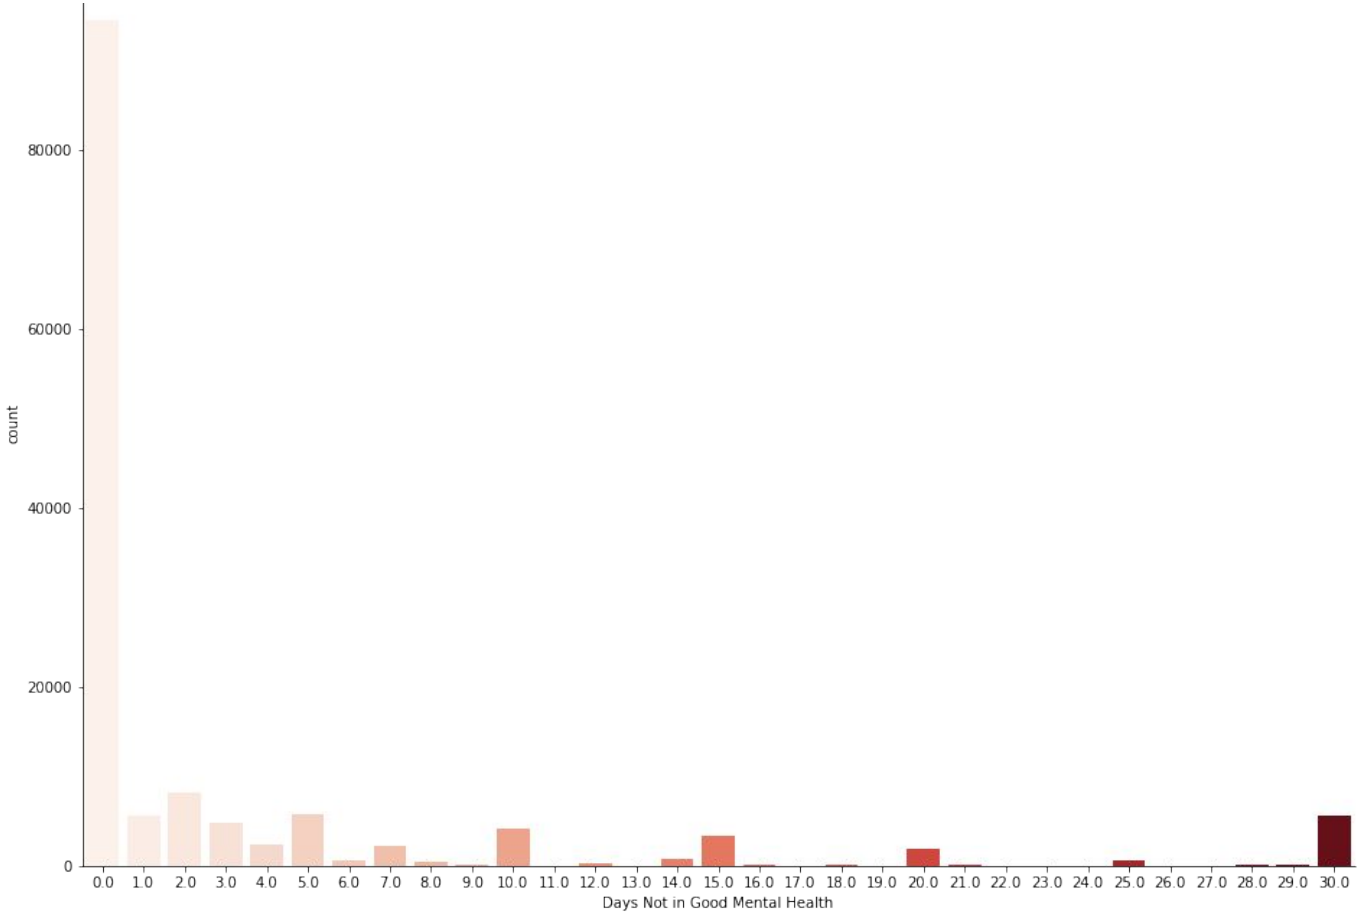

An equal number of respondents were diagnosed with skin cancer as all other forms of cancer

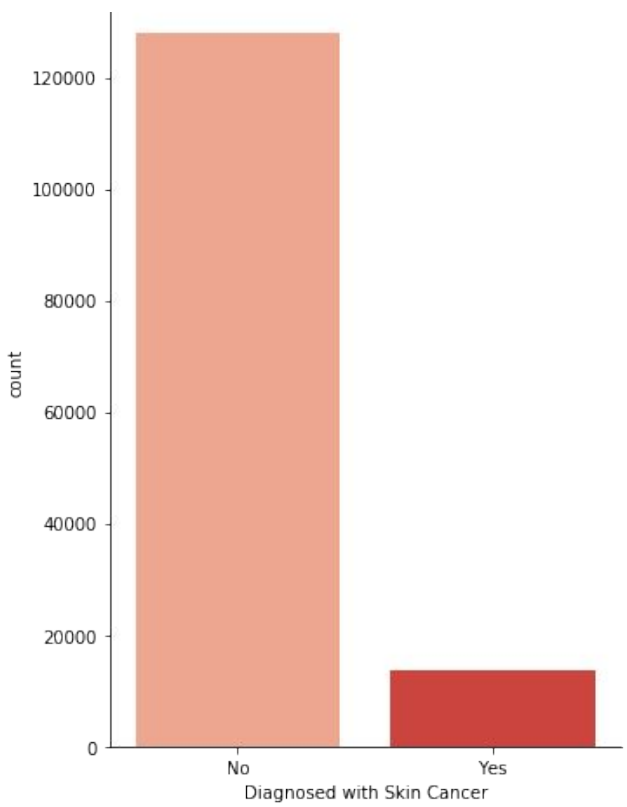

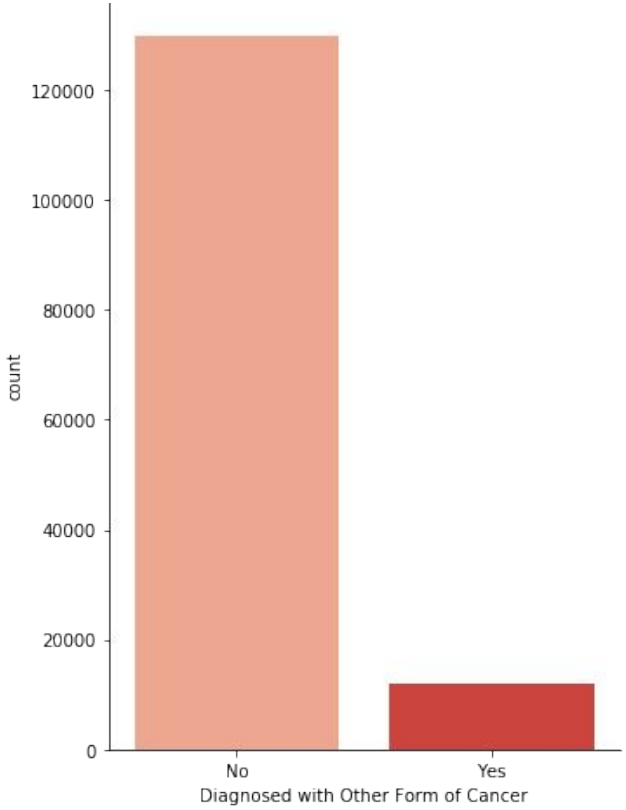

Just 6% of respondents were heavy drinkers

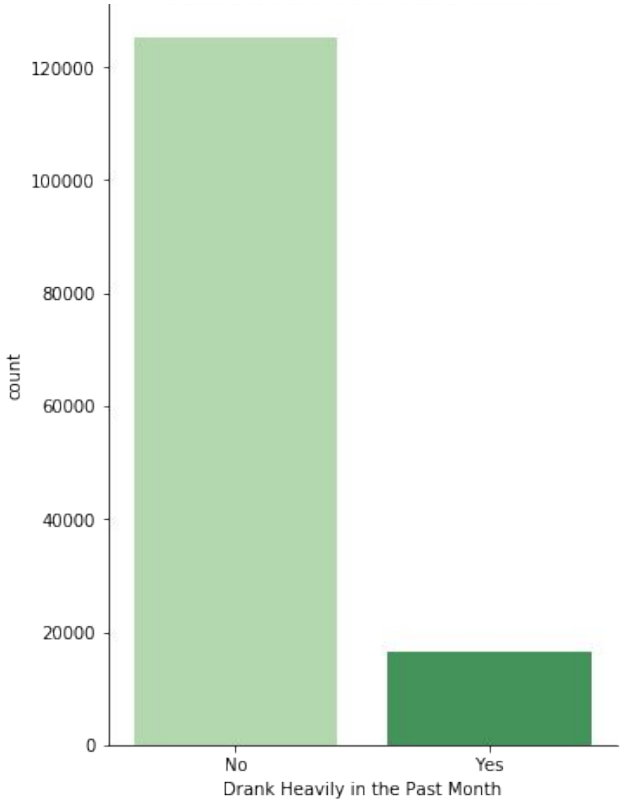

Most respondents always wore their seatbelt

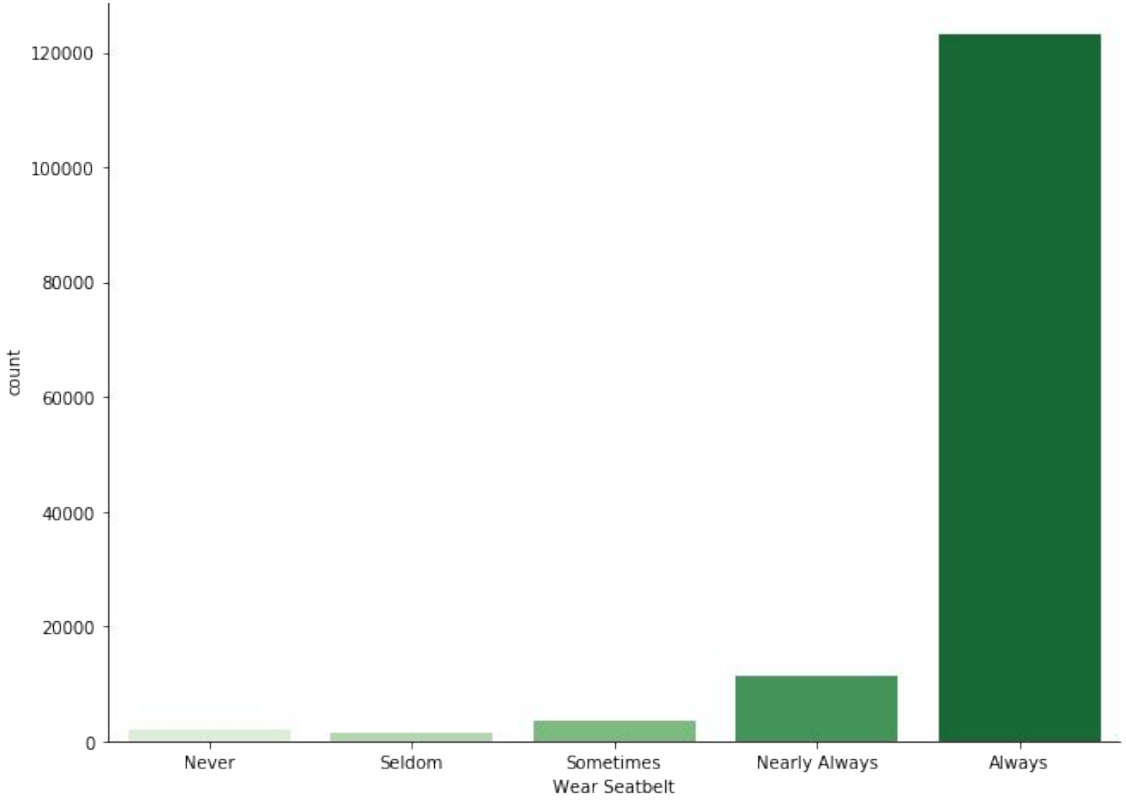

10% of respondents were former smokers, while 4% smoked every day

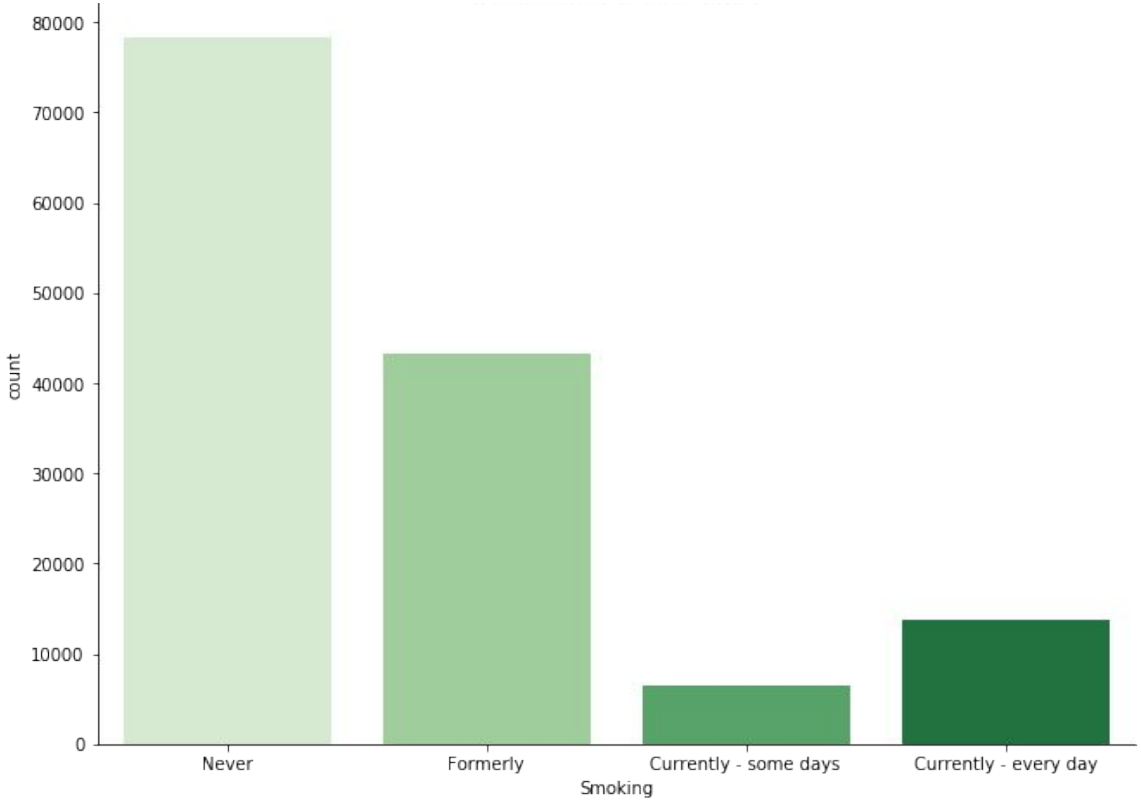

Less than 1% of respondents were considered at risk for HIV

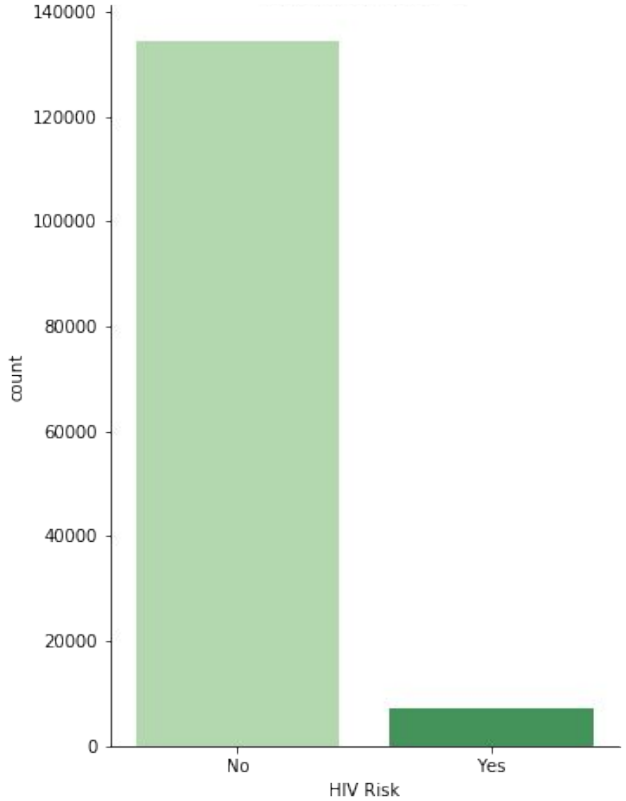

Methods

- Feature selection
- Clustering
- Interpretation

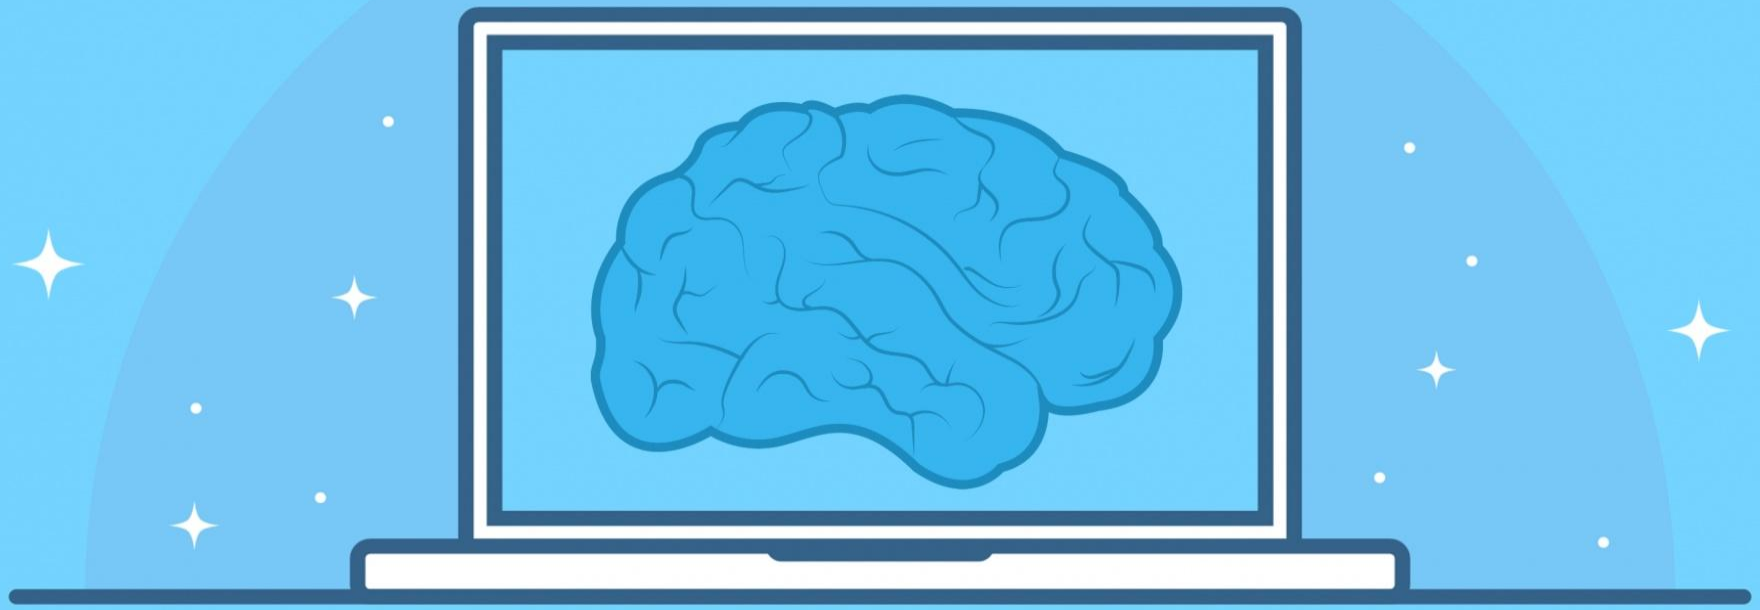

Feature Selection - Multicollinearity Heatmap

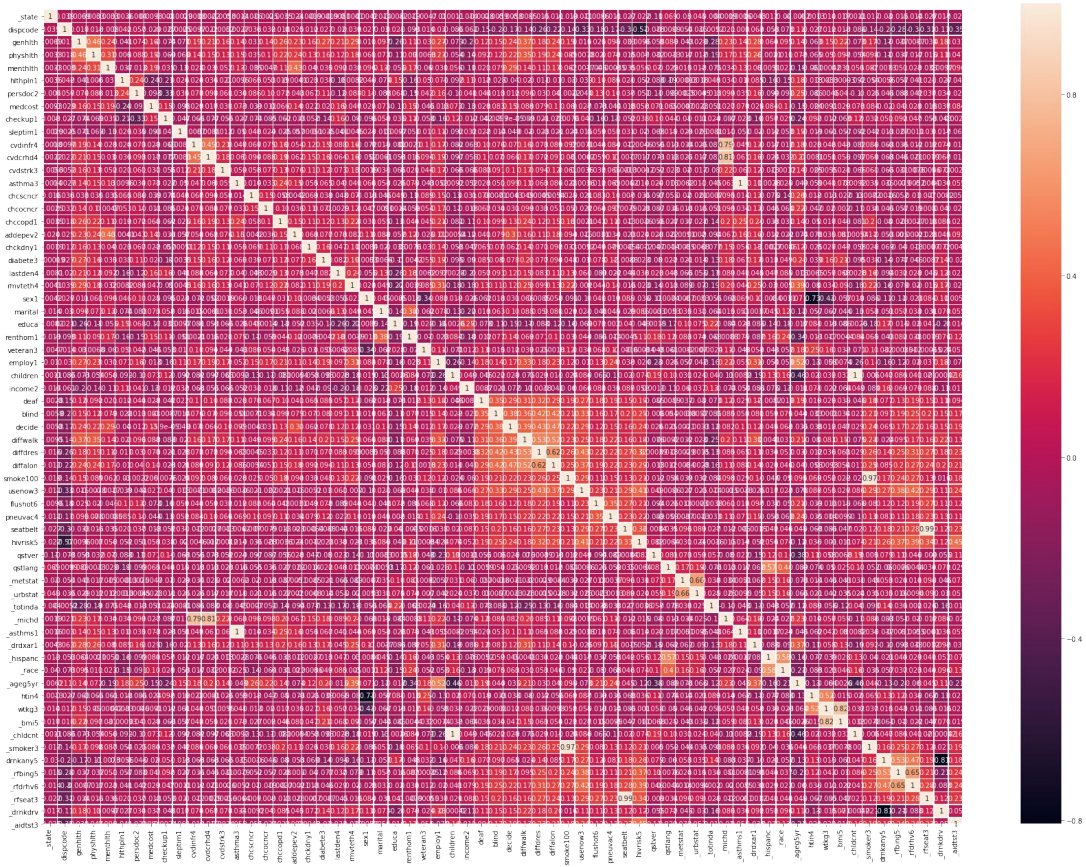

Feature Selection - Variance Inflation Factor

Result: Removed total of 13 variables

Total: 55 features & 437,436 observations

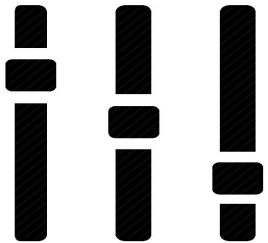

DBSCAN

Result: Too much noise - 99.99%

Best DBSCAN: .18 eps, 300 min_samples
→ 7 clusters

Silhouette Score: -.45

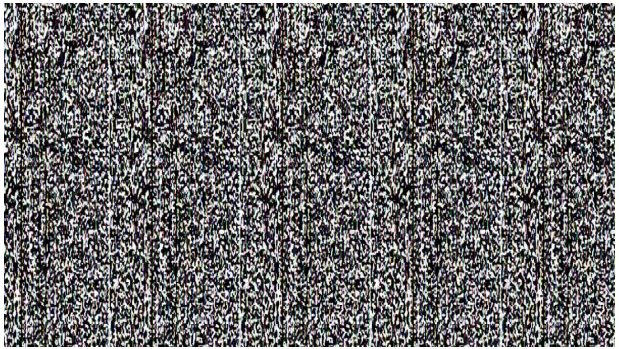

K-Modes

K=7

Score: .04

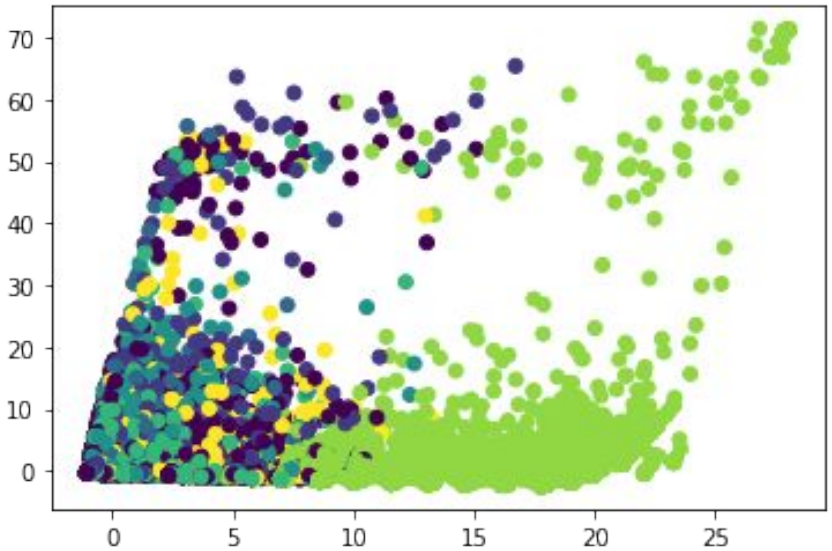

K-Means

K=7

Score: .25

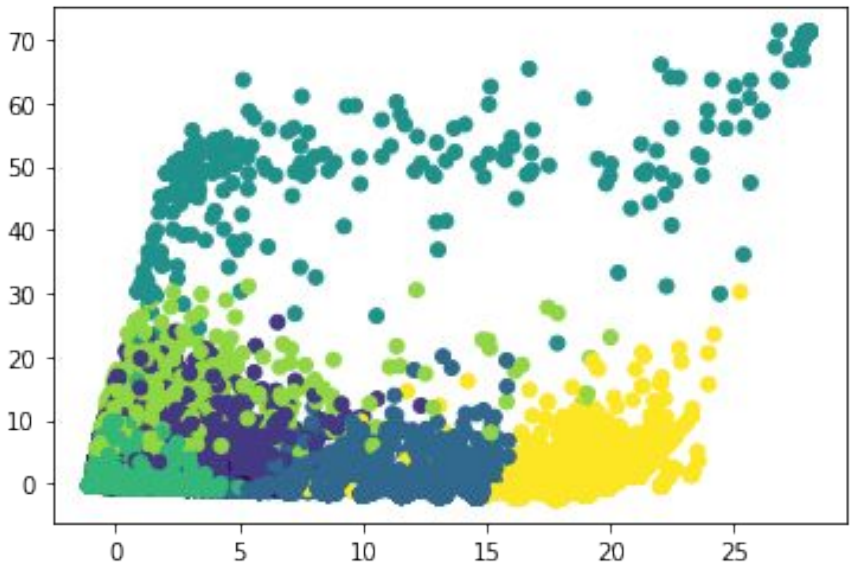

Cluster Breakdown

Cluster Number	Count
0	131,595
4	121,006
1	63,451
3	60,640
5	40,294
2	11,720
6	8,730

Clusters 3 & 6 were more likely to be retired or unable to work

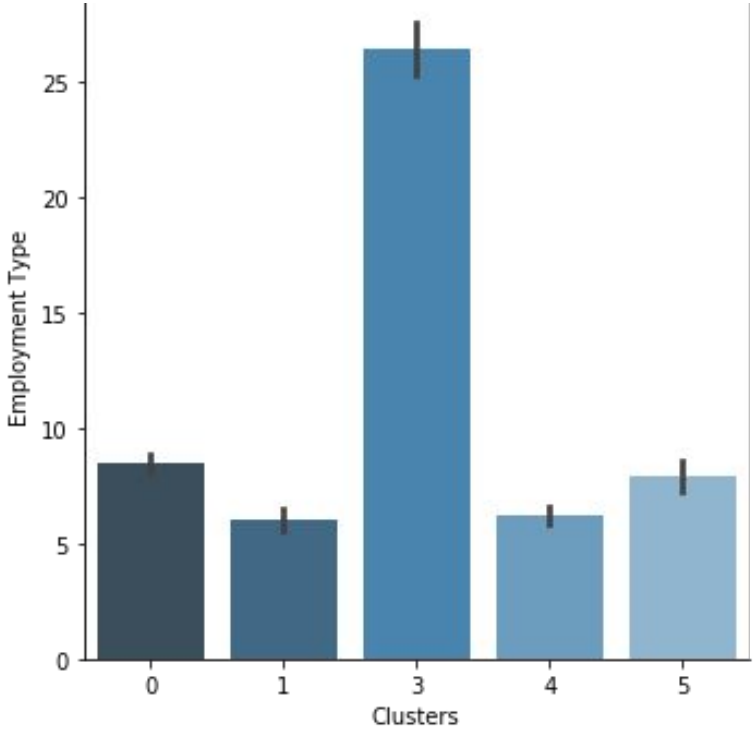

Cluster 3's average sleep time was highest

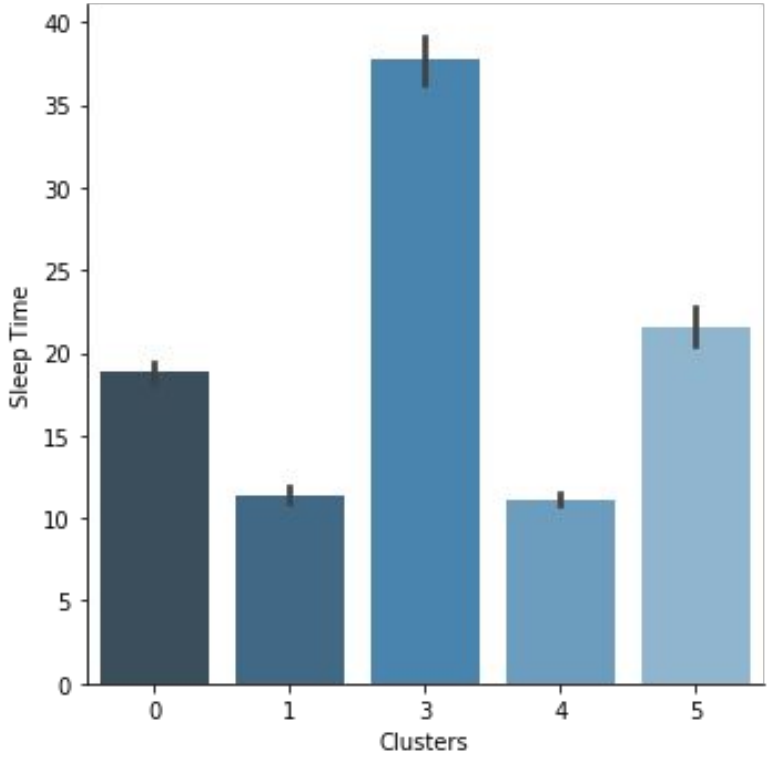

Clusters 3, 1, & 4 were more likely to have drunk alcohol recently

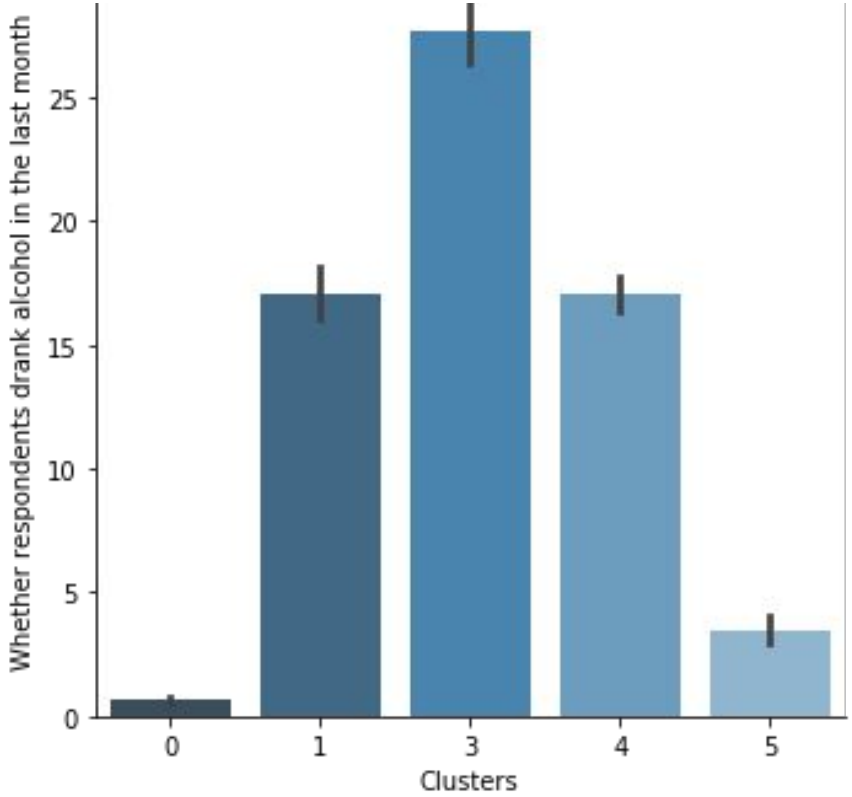

Clusters 3 & 0 were older on average

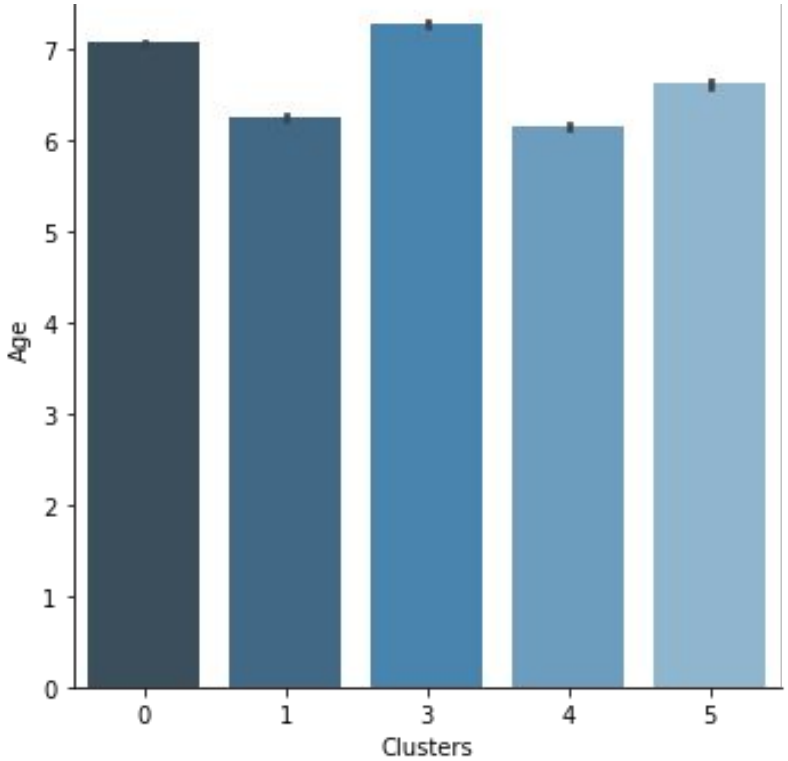

Limitations

- Slowness of running the algorithms
- Could have tried a few difference values for eps, if there was more time
- Categorical data not the best to cluster with

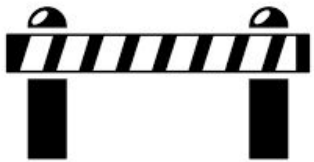

Conclusions & Future Research Possibilities

- Possibility that the dataset was clustered based on life stages of respondents
- Potential other aspects at play too, such as affluence and lifestyle
- Would be nice to learn more about the clusters and plot them alongside more features
- Would have been good idea to try HDBSCAN